

Deep Learning for Sentiment Analysis: Case Studies

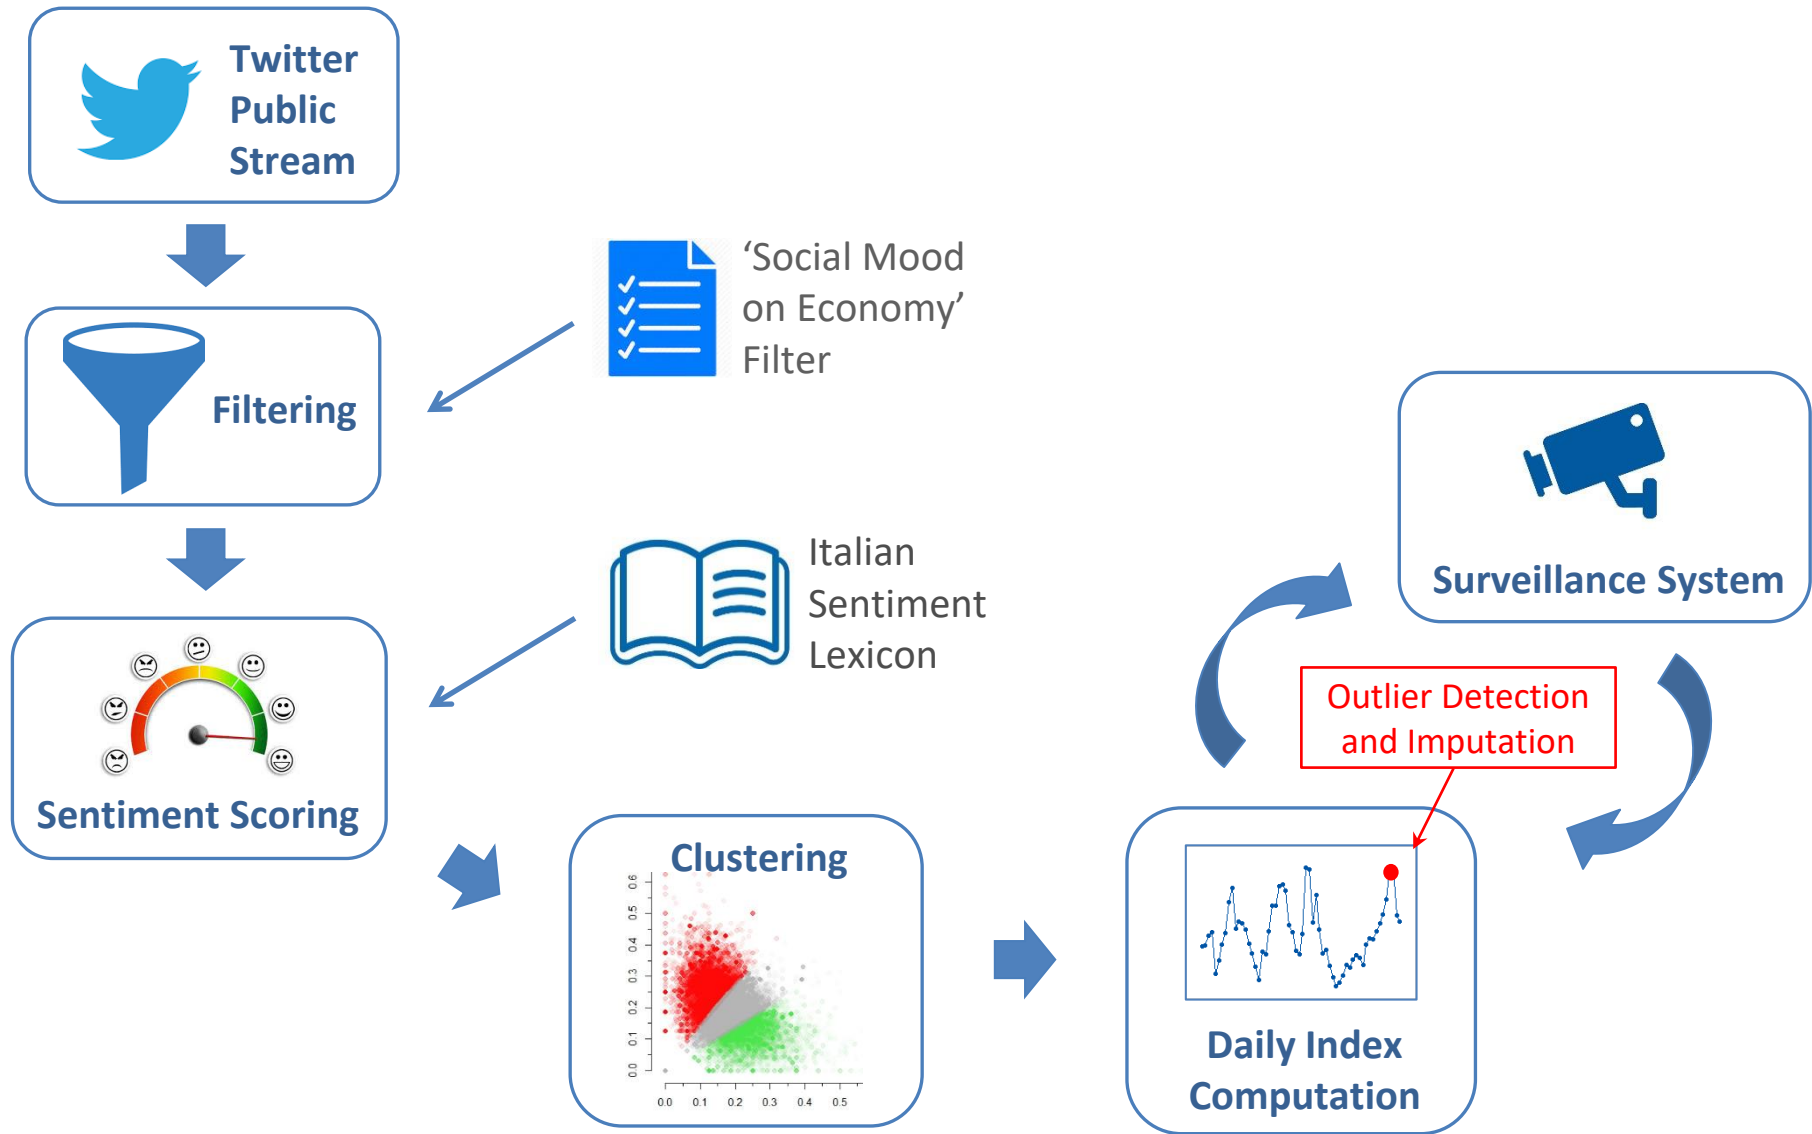
Francesco Pugliese, PhD

Email Francesco Pugliese : neural1977@gmail.com

Case Study: Sentiment Analysis for Economy

- Nowadays more and more people are using Social Media platforms to find out news, to express their feelings and to share or debate opinions about virtually every possible topic
 - The interest towards Social Media as a means for “measuring” public’s mood is **still growing**
- We are investigating whether social media messages may be successfully exploited to develop **domain-specific sentiment indices**. The aim is to assess the **Italian mood about specific topics or aspects of life**, e.g.
 - the economic situation, the European Union, the migrants’ phenomenon, the terrorist threat, and so on
- These new indices would enable **high-frequency** (e.g. **daily**) measures of the Italian sentiment about phenomena which are of **interest in Official Statistics**
- The hope is that such indices could either improve the performance of forecasting models, or enrich existing statistical products (e.g. the BES), or even be disseminated as **new statistical outputs in their own right**

Processing Pipeline at a Glance



Data Collection and Storage

- **Data Collection Technique**
 - We exploit **Twitter's Streaming API** to get low latency access to Twitter's firehose and collect samples of **public tweets**
- **Target Population**
 - Public tweets whose text matches **at least one keyword** belonging to the filter
- **Sampling Design**
 - The sampling algorithm is a **black box**, as it is entirely controlled by Twitter's Streaming API. **At most a 1%** of all the tweets produced on Twitter at a given time can be sampled
- **Data Format**
 - Twitter's Streaming API returns data in **JSON** format
- **Data Staging and Storage**
 - We temporarily store gathered JSON data as text files inside a **staging area** residing on a server. Then we periodically load bunches of tweets into an **Oracle DB** (and remove the corresponding files from the staging area)

Text Processing and Sentiment Analysis

- **Process Granularity**
 - To compute **daily index values**, we process all the tweets collected in a **single day** as a **single block**
- **Input Data**
 - We **only** analyze **the textual content** of the tweets. (No information about users is ever accessed: the index only uses *unlinked anonymized* data)
- **Text Cleaning and Normalization**
 - We perform standard NLP **pre-processing** steps: (i) convert to lowercase, (ii) tokenize running text into words, (iii) apply basic orthographic repairs, (iv) remove URLs, (v) remove non-alphabetic characters (e.g. '#' or '@'), (vi) remove stop words, (vi) *if needed*, **stem** words to get rid of inflected forms
- **Sentiment Analysis Approach**
 - To classify tweets as Positive, Negative or Neutral we chose to adopt an **unsupervised, lexicon-based approach**
 - We discarded supervised, Machine Learning approaches because we were **unable to find** large, high quality **training sets** of human-labeled tweets **in Italian**

Sentiment Scores: The Lexicon


- Our Sentiment Analysis application involves two sequential steps
 - 1) Calculate **sentiment scores** for each tweet
 - 2) Use these **sentiment scores** to **cluster** tweets into three mutually exclusive classes: Positive (P), Negative (N) and Neutral (U)
- To attach sentiment scores to a tweet we leverage an **Italian Sentiment Lexicon**, namely a vocabulary whose lemmas are associated to **pre-computed positive** and **negative** sentiment scores
- Currently we are using the **Sentix** lexicon [Basile and Nissim 2013]
- Since it aligns several existing, **independent** lexical resources (*WordNet*, *MultiWordNet*, *BabelNet*, *SentiWordNet*) Sentix contains many **duplicated** lemmas
 - ~75'000 lemmas overall, only ~42'000 **unique**
 - ➡ To ensure unambiguous and reproducible results we **de-duplicated** Sentix by **averaging** atomic sentiment scores of duplicated lemmas

The Sentiment Space

- In Sentix, positive (p) and negative (n) **sentiment scores** of lemmas are **constrained** as follows:

$$\begin{cases} p \in [0, 1] \\ n \in [0, 1] \\ p + n \leq 1 \end{cases}$$

- Therefore Sentix maps lemmas to **points** belonging to the **sentiment triangle**:



- From (p, n) coordinates we can pass to **polar** coordinates (i, θ) and derive two **additional** sentiment scores:

✓ **Polarity** $\omega = 1 - 4\theta/\pi$ $\omega \in [-1, 1]$

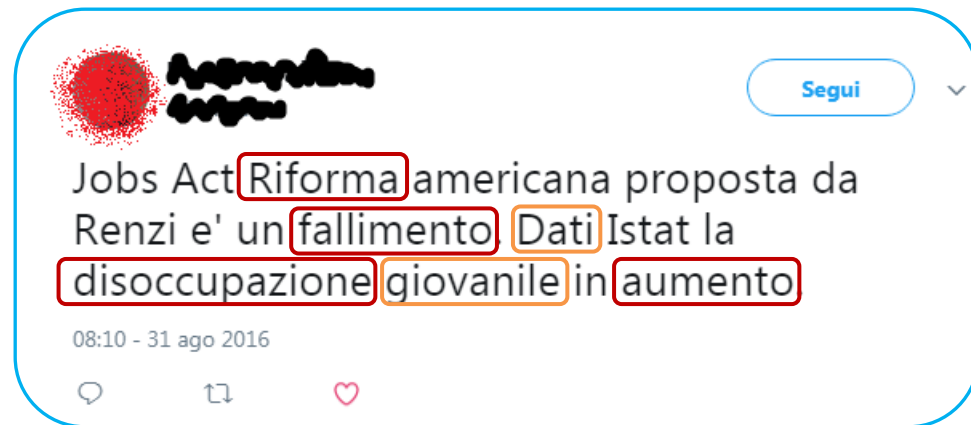

✓ **Intensity** $i = \sqrt{p^2 + n^2}$ $i \in [0, 1]$

- This way Sentix lemmas are mapped to a **4D sentiment space**


lemma	pos	neg	polarity	intensity
caldo	0.25	0.125	0.41	0.28
freddo	0.047	0.297	-0.8	0.3

- ➔ To enable clustering, **tweets too** must be mapped to this 4D space

From Word-level to Tweet-level Sentiment Scores



word	pos	neg	polarity	intensity
riforma	0.125	0	1	0.125
fallimento	0.021	0.375	-0.929	0.376
dat	0.063	0.104	-0.312	0.121
disoccupazione	0	0.625	-1	0.625
giovani	0.125	0	1	0.125
aumento	0.208	0.083	0.516	0.224




pos	neg	polarity	intensity
0.090	0.198	-0.455	0.218

Clustering and Calculation of the Index

- Once sentiment scores (p, n, ω, i) are available for all the tweets of a daily block...
- ...we use **K-means** to **cluster** them into Positive, Negative and Neutral tweets
 - ✓ to lower the risk of finding a local optimum, we run it 100 times with random starts and pick the best solution
- Lastly we compute the **daily index value** (S), which depends on the distribution of tweets within the Positive, Neutral and Negative classes

$$S = \bar{\omega}_i = \frac{\sum_t i_t \omega_t}{\sum_t i_t} = \frac{\sum_{t \in P} i_t \omega_t + \sum_{t \in N} i_t \omega_t}{\sum_t i_t}$$

where $\omega_t \stackrel{\text{def}}{=} 0 \quad \forall t \in \text{Neutral}$



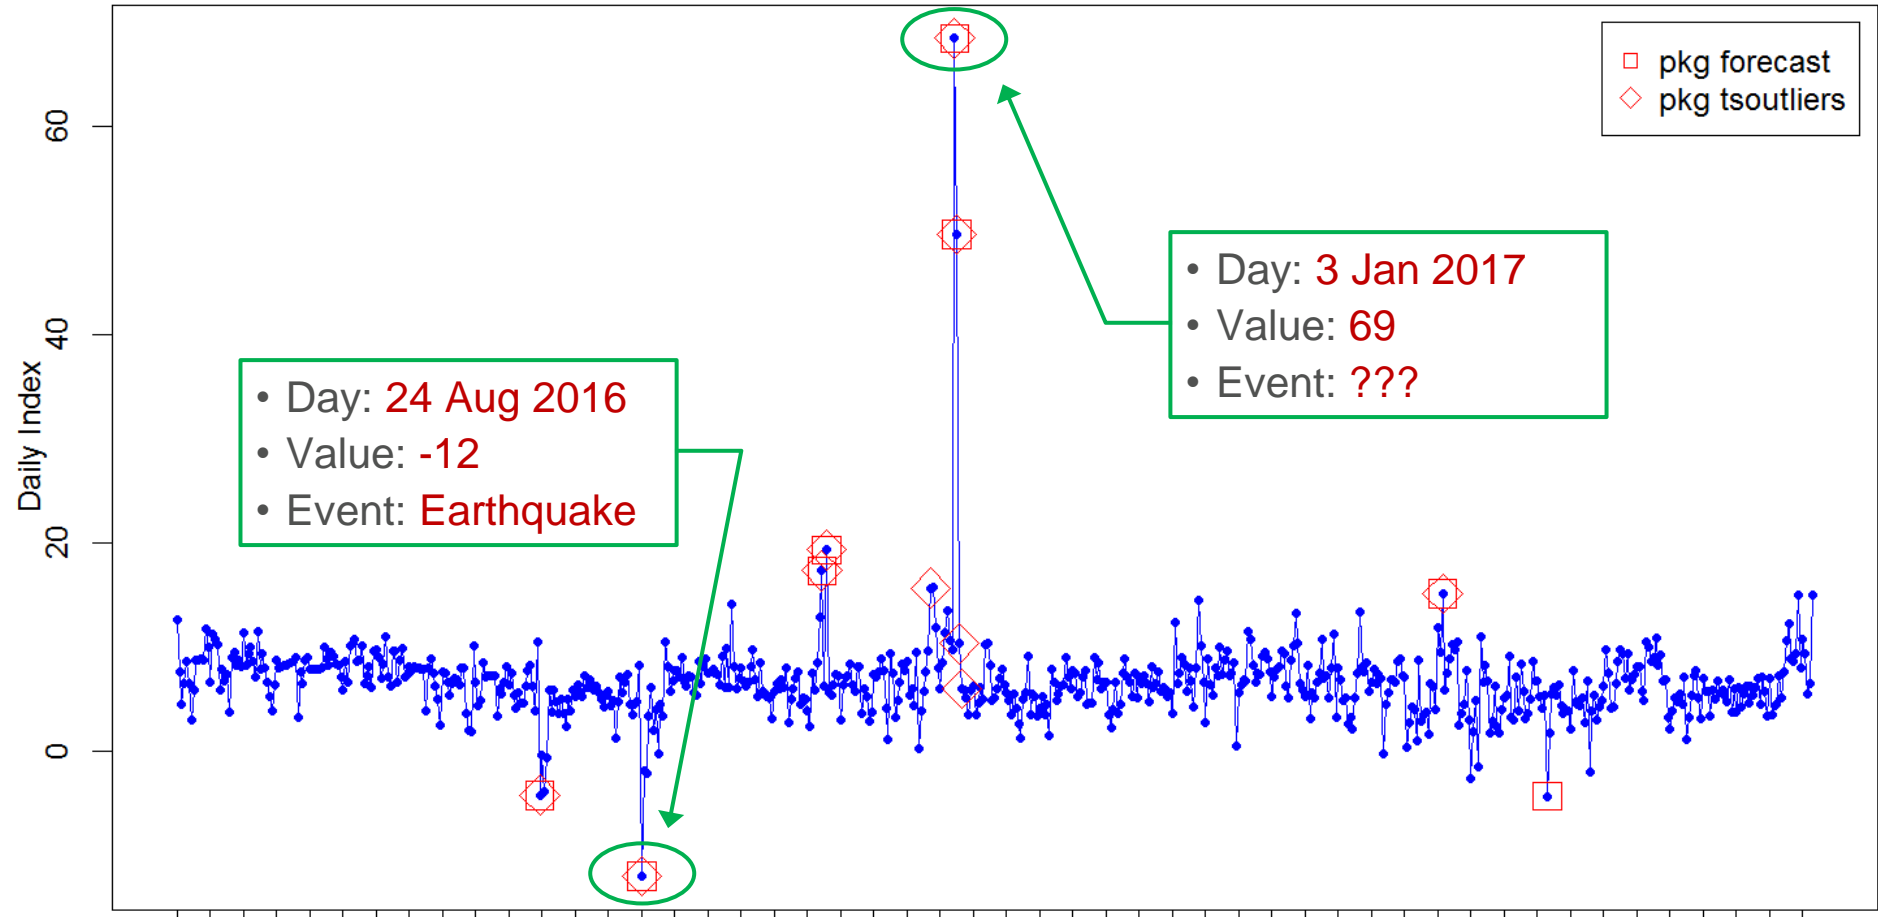
- This index can be seen as the **average of polarity** (ω) **weighted by intensity** (i), provided we treat *Neutral* tweets as if their polarity were zero. Compared to traditional alternatives:
 - ✓ It is more **resilient** to tweets' **misclassification**
 - ✓ It **reduces** day-to-day **volatility**

Monitoring and Validation

- *No filter is perfect!* Thus we devoted special care to make the index **robust** against possible **contaminations by off-topic tweets** that might pass the filter
- We developed a **surveillance system**, which periodically searches for **anomalous values** in the daily time series by means of **two** independent and complementary **outlier detection routines**
 - Daily values detected as **potential outliers** cause the system to generate a set of **automated diagnostic reports**
 - These are then sent to **human reviewers** in charge of deciding whether the detected values are actually proper data points, or instead **truly anomalous**
- Truly anomalous data typically arise when an **off-topic** tweet that happened to pass the filter becomes **“viral”** on Twitter
 - Being re-tweeted and quoted thousands of times in a day, viral tweets may have an **unduly impact** on the daily index and introduce **bias**
- All the daily index values classified as **truly anomalous** are eventually **imputed** via nearest-neighbor interpolation

Anomalous Values: One Example

Check for Suspect Outliers




prima
fedefederossi
ben
spesa




Volume Burst in 2018 Post-Election Crisis

Social Mood on Economy - Daily Volume




Valleys: Disasters and Terrorism

Social Mood on Economy - Daily Index and Moving Averages




Valleys: Welfare and Economy

Social Mood on Economy - Daily Index and Moving Averages




Peaks: Welfare and Economy

Social Mood on Economy - Daily Index and Moving Averages




Peaks: Holidays

Social Mood on Economy - Daily Index and Moving Averages



An Interesting Dynamic

Social Mood on Economy - Daily Index and Moving Averages



Sentiment Analysis (*Ain, et al. 2017*)

- **Sentiments** of users that are expressed on the web has great influence on the readers, product vendors and politicians.
- **Sentiment Analysis** refers to text organization for the classification of mind-set or feelings in different manners such as negative, positive, favorable, unfavorable, thumbs up, thumbs down, etc. Thanks to DL, the SA can be visual as well.



Discovering people opinions, emotions and feelings about
a product or service

Sentiment Analysis Using Deep Learning Techniques: A Review

Qurat Tul Ain*, Mubashir Ali*, Amna Riaz[†], Amna Noureen[‡], Muhammad Kamran[‡], Babar Hayat* and A. Rehman*

*Department of Computer Science and Information Technology,
The University of Lahore, Gujrat, Pakistan

[†]Department of Information and Technology,
University of Gujrat, Gujrat, Pakistan

[‡]Department of Computer Engineering,
EME, NUST, Islamabad, Pakistan

a) Machine learning based techniques: This type of techniques are implemented by extracting the sentences and aspect levels. The features consist of Parts of Speech (POS) tags, n-grams, bi-grams, uni-grams and bag-of-words. Machine learning contains three flavors at sentence and aspect, i.e., Nave Bayes, Support Vector Machine (SVM) and Maximum Entropy.


b) Lexicon based or corpus based techniques: These techniques are based on decision trees such as k-Nearest Neighbors (k-NN), Conditional Random Field (CRF), Hidden Markov Model (HMM), Single Dimensional Classification (SDC) and Sequential Minimal Optimization (SMO), related to methodologies of sentiment classification.

Sentiment Analysis with Feedback

Stockle [start page](#)




Apple Inc. **AAPL** 116.30 (+0.25%)



ADBE **ADBE** 0.0 (0.0%)




eBay Inc. **EBAY** 31.46 (-0.49%)



GOOGL **GOOGL** 0.0 (0.0%)



Microsoft Corporation **MSFT** 57.19 (-0.85%)




Yahoo! Inc. **YHOO** 42.68 (-1.24%)



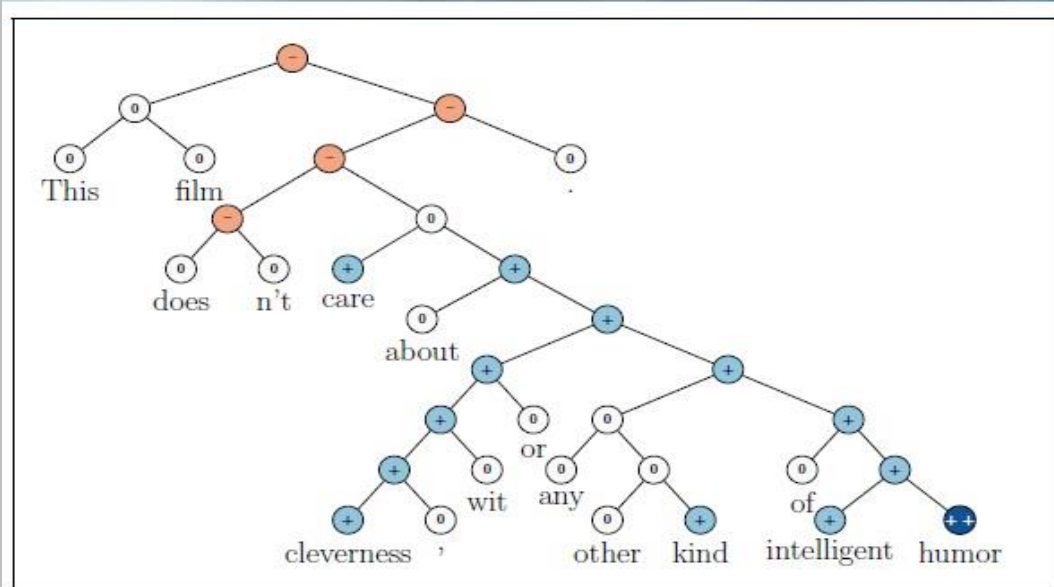
Recursive Neural Networks (RecursiveNN) (Socher, R., et al., 2011b)

- These models are recursive auto-encoders which learn semantic vector representations of phrases. Word indices (orange) are first mapped into a semantic vector space (blue).
- Then they are recursively merged by the same auto-encoder network into a fixed length sentence representation. The vectors at each node are used as features to predict a distribution over text labels.





Recursive Neural Tensor Networks (RNTN) (Socher, R., et al. 2013)

- The Stanford Sentiment Treebank is the first corpus with fully labeled parse trees that allows for a complete analysis of the compositional effects of sentiment in language.
- RNTNs compute parent vectors in a bottom up fashion using a compositionality function and use node vectors as features for a classifier at that node.



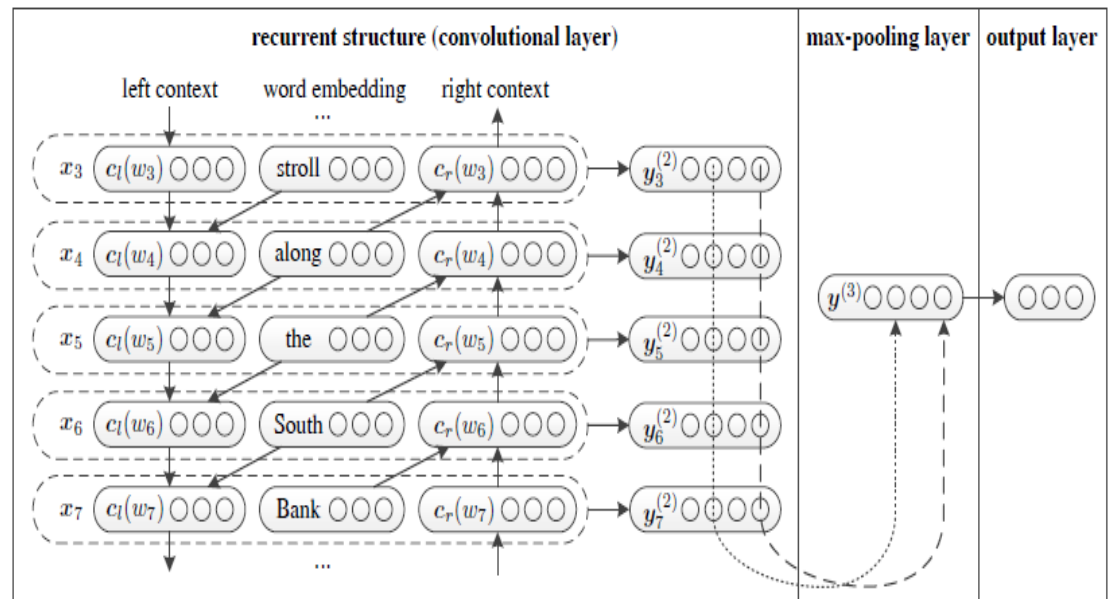
RNTN – Upside and Downside

- RNTNs are very efficient in terms of constructing sentence representations.
- RNTNs capture the semantics of a sentence via a tree structure. Its performance heavily depends on the performance of the textual tree construction.
- Constructing such a textual tree exhibits a time complexity of at least $O(n^2)$, where n is the length of the text.
- RNTNs are unsuitable for modeling long sentences or documents.



Recurrent Convolutional Neural Networks (RCNN) (*Lai, S., et al. 2015*)

- They adopt a recurrent structure to **capture contextual information** as far as possible when learning word representations, which may introduce considerably **less noise compared** to traditional window-based neural networks.
- The **bi-directional recurrent structure** of RCNNs.
- **RCNNs** exhibit a time complexity of $O(n)$



RCNN Equations

- RCNNs exhibit a **time complexity of $O(n)$** , which is linearly correlated with the length of the text length.

$$c_l(w_i) = f(W^{(l)}c_l(w_{i-1}) + W^{(sl)}e(w_{i-1})) \quad (1)$$

$$c_r(w_i) = f(W^{(r)}c_r(w_{i+1}) + W^{(sr)}e(w_{i+1})) \quad (2)$$

- 7 equations** defining all the Neural Network topology

$$x_i = [c_l(w_i); e(w_i); c_r(w_i)] \quad (3)$$

$$y_i^{(2)} = \tanh(W^{(2)}x_i + b^{(2)}) \quad (4)$$

$$y^{(3)} = \max_{i=1}^n y_i^{(2)} \quad (5)$$

- Input length** can be variable

$$y^{(4)} = W^{(4)}y^{(3)} + b^{(4)} \quad (6)$$

$$p_i = \frac{\exp(y_i^{(4)})}{\sum_{k=1}^n \exp(y_k^{(4)})} \quad (7)$$

RCNN: Feature Extraction

- RCNNs employ a max-pooling layer that automatically judges which words play key roles in text classification to capture the key components in texts.
- The most important words are the information most frequently selected in the max-pooling layer.
- Contrary to the most positive and most negative phrases in RNTN, RCNN does not rely on a syntactic parser, therefore, the presented n-grams are not typically “phrases”.

RCNN

P	well <i>worth</i> the; a <i>wonderful</i> movie; even <i>stinging</i> at; and <i>invigorating</i> film; and <i>ingenious</i> entertainment; and <i>enjoy</i> .; 's <i>sweetest</i> movie
N	A <i>dreadful</i> live-action; Extremely <i>boring</i> .; is <i>n't</i> a; 's <i>painful</i> .; Extremely <i>dumb</i> .; an <i>awfully</i> derivative; 's <i>weaker</i> than; incredibly <i>dull</i> .; very <i>bad</i> sign;

RNTN

P	an amazing performance; most visually stunning; wonderful all-ages triumph; a wonderful movie
N	for worst movie; A lousy movie; a complete failure; most painfully marginal; very bad sign

RCNN applied to social networks

Sentiment Analysis: Twitter

- RCNNs achieve 85% of accuracy on a 1.6 mln tweets training set (800k positive and 800k negative) for the task of Sentiment Analysis. This result is «state-of-art» in twitter sentiment classification.
- We can extract most significant keywords and summarize

```
Tweet: Played with an android google phone. The slide out screen scares me I would break that fucker so fast. Still prefer my iPhone.  
- Sentiment: -0.98 - -1
```

```
Keywords: prefer, still, me, fucker
```

```
Tweet: US planning to resume the military tribunals at Guantanamo Bay... only this time those on trial will be AIG execs and Chrysler debt holders  
- Sentiment: -0.51 - -1
```

```
Keywords: only, holders, aig, debt
```

```
Tweet: omg so bored & my tattoooooos are so itchy!! help! aha =>  
- Sentiment: -0.99 - -1
```

```
Keywords: aha, itchy, bored, help
```

```
Tweet: I'm itchy and miserable!  
- Sentiment: -1.00 - -1
```

```
Keywords: miserable, itchy, and
```

```
Tweet: @sekseemess no. I'm not itchy for now. Maybe later, lol.  
- Sentiment: 1.00 - +1
```

```
Keywords: lol, later, itchy, maybe
```

RCNN applied to Extractive Text Summarization

- **Best keywords lead to best contextes ---> Summarization**

```
Tweet 29: "Già avete letto 136 pagine del piano scuola? #Fenomeni #labuonascuola"
```

```
Sentiment: -0.95 - -1
```

```
Keywords: pagine, avete, fenomeni, piano
```

```
Tweet 30: "\"Per l'aternanza #scuola #lavoro bisogna passare da 11a 100milioni di euro\" #labuonascuola http://t.co/zGAzknI8rv"
```

```
Sentiment: -0.81 - -1
```

```
Keywords: euro, t, scuola, lavoro
```

```
Most significant keywords driving the sentiment decision:
```

```
Eccolo
```

```
Siamo
```

```
Scuola
```

```
Giuste
```

```
Escluso
```

```
Most significant sentences driving the sentiment decision:
```

```
...cambierà solo se noi metteremo al centro...
```

```
...solo se noi metteremo al centro la...
```

```
...più grande spettacolo mai visto passodopopasso scuola...
```

```
...mai visto passodopopasso scuola labuonascuola...
```

```
...nessuno si senta escluso la buona scuola...
```

Recurrent Neural Networks are able to understand negations and other things

- Thanks to **word embeddings** semantics **RNNs** can recognize **negations**, and complex forms of language utterances.

Tweet: This is a bad thing
- Sentiment: -0.72 - -1

Keywords: bad, thing, a, is

Tweet: This is not a bad thing
- Sentiment: 0.46 - +1

Keywords: not, thing, bad, a

Tweet: This is a positive thing
- Sentiment: 0.94 - +1

Keywords: positive, thing, a, is

Tweet: This is a very positive thing
- Sentiment: 0.91 - +1

Keywords: positive, very, thing, a

Tweet: I like Renzi politics
- Sentiment: 0.70 - +1

Keywords: like, renzi, politics, i

Tweet: I don't agree with Renzi Politics
- Sentiment: 0.16 - 0

Keywords: don't, agree, politics, renzi

Tweet: Renzi did a wrong international Politics
- Sentiment: -0.34 - -1

Keywords: wrong, did, renzi, international

Tweet: Renzi did a very good international Politics
- Sentiment: 0.74 - +1

Keywords: did, renzi, good, very

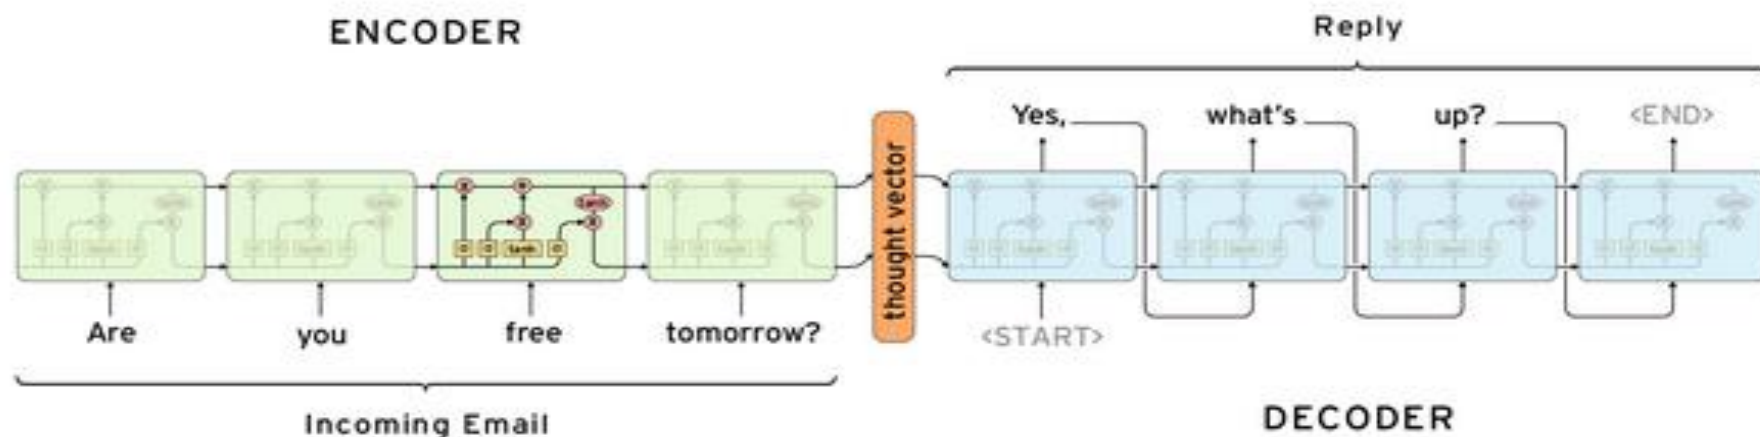
Tweet: Istat is a very good Institute of research
- Sentiment: 0.84 - +1

Keywords: good, very, research, istat

Tweet: Istat is not a good Institute of research - Sentiment: -0.78 - -1

Keywords: not, research, istat, institute

Neural Conversational Models (Vinyals, & Le., 2015).



Conversation model – chatbot?

- Training on a set of conversations. The input sequence can be the concatenation of what has been conversed so far (the context), and the output sequence is the reply.

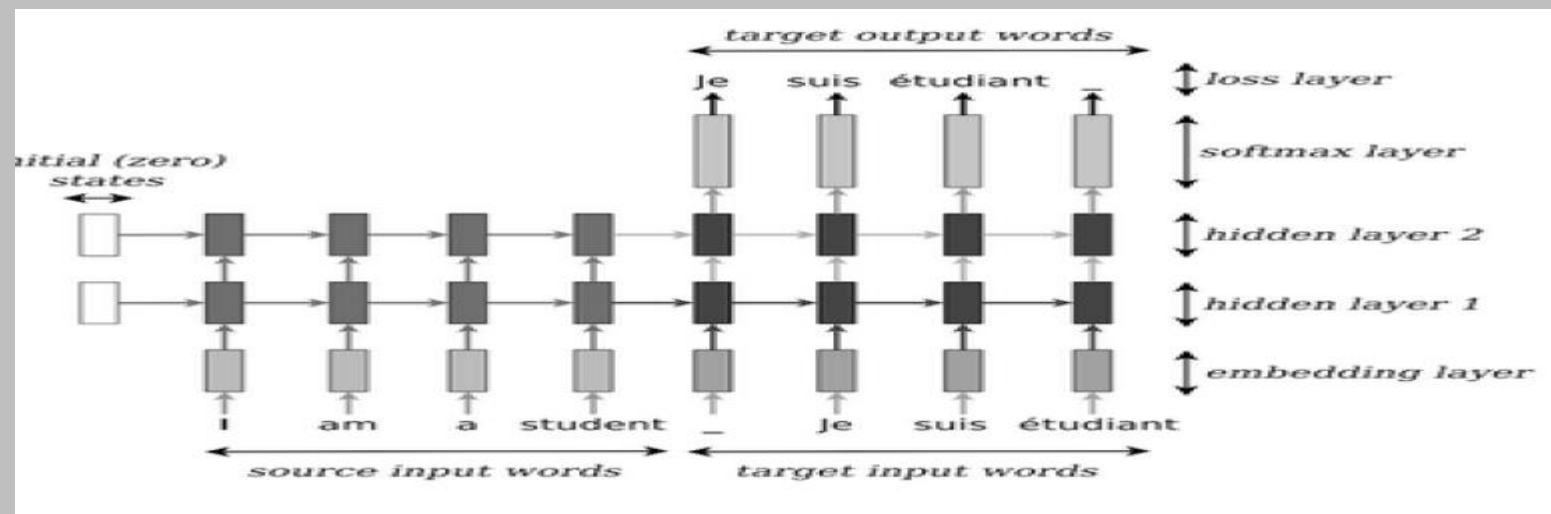
Conversation 4: Philosophical Q&A

Machine: *is it a windows machine ?*
Machine: *we have linux on your machine .*
Machine: *is it windows xp or windows 7 ?*
Human: *i am using linux*
Machine: *let me take a remote of your machine and check if issue has been pushed properly .*
Human: *ok*
Machine: *please open terminal and run the below commands in one computer :*
Machine: *sudo apt - get update*
Machine: *sudo apt - get upgrade*
Machine: *sudo <COMMAND> - updater*
Human: *done*
Machine: *let me try connecting now*
Human: *ok*

Human: *what is the purpose of life ?*
Machine: *to serve the greater good .*
Human: *what is the purpose of living ?*
Machine: *to live forever .*
Human: *what is the purpose of existence ?*
Machine: *to find out what happens when we get to the planet earth .*
Human: *where are you now ?*
Machine: *i 'm in the middle of nowhere .*
Human: *what is the purpose of dying ?*
Machine: *to have a life .*
Human: *what is the purpose of being intelligent ?*
Machine: *to find out what it is .*
Human: *what is the purpose of emotions ?*
Machine: *i don 't know .*

Neural Machine Translation (Bahdanau, et., al., 2014).

- Neural machine translation (NMT) is an approach to “machine translation” that uses large ANN to predict the likelihood of a sequence of words, typically modeling entire sentences in a single integrated model (Bahdanau et al., 2014; Luong et Manning, 2016).



Multilingual Sentiment Analysis (Testi, M., Mercuri, A., & Pugliese, F. Multilingual Sentiment Analysis. *JADT'18*, 780.)

- During the training stage, the RCNN achieves 84% of accuracy on a validation set (selected at the 20% of the original dataset). On a test set of 380 tweets (provided by Semeval), the model returns around 82% of accuracy on positive tweets and 78% of accuracy on negatives, with an approximative 80% overall on a mixed tweets set.
- During the training we determined 3.2 millions of keywords, namely 2 for each tweet, the most important and the second in order of significance.



Contextual Translations Web-sites

python - Con 190424288-3c numpy.load.py python - How Migrants, Fran (54) ESTETICA Francesco Pu Jadt 2018 - Intern Neural mach politica - T X

context.reverso.net/traduzione/italiano-inglese/politica 90% Cerca

Reverso Context TRADUZIONE | DIZIONARIO | CORRETTORE ORTOGRAFICO | CONIUGAZIONE | GRAMMATICA

politica Italiano Inglese

Forse intendi: politico

Traduzione di "politica" in inglese Vedi anche: politica estera politica agricola politica europea polit...

Sostantivo **Aggettivo**

policy politics policy-making politician behaviour policymaking policymaker political policies politically

affairs strategy stance agenda EU

Dovevamo discutere un'importante iniziativa politica. He and I were supposed to discuss a major policy initiative.

La relatrice desidera incoraggiare questa politica più "interventista". The rapporteur would like to encourage this more "activist" policy.

Questo influenzerebbe la politica interna di tutti i Kazon. That would clearly affect the internal politics of all the Kazon.

Ho sempre voluto entrare in politica. I have always wanted to get into politics.

Bene avere più donne impegnate in politica. It's good to get more women involved in politics.

Iscriviti a Reverso

Pubblicità

PrestitiOnline.it
Più scelta, più risparmio

Richiedi il tuo Prestito

PrestitiOnline.it: prestiti personali, auto, cessione del quinto e consolidamento debiti.

Pubblicità

FLUKE

Introduzione al nuovo multimetro Fluke 279 FC

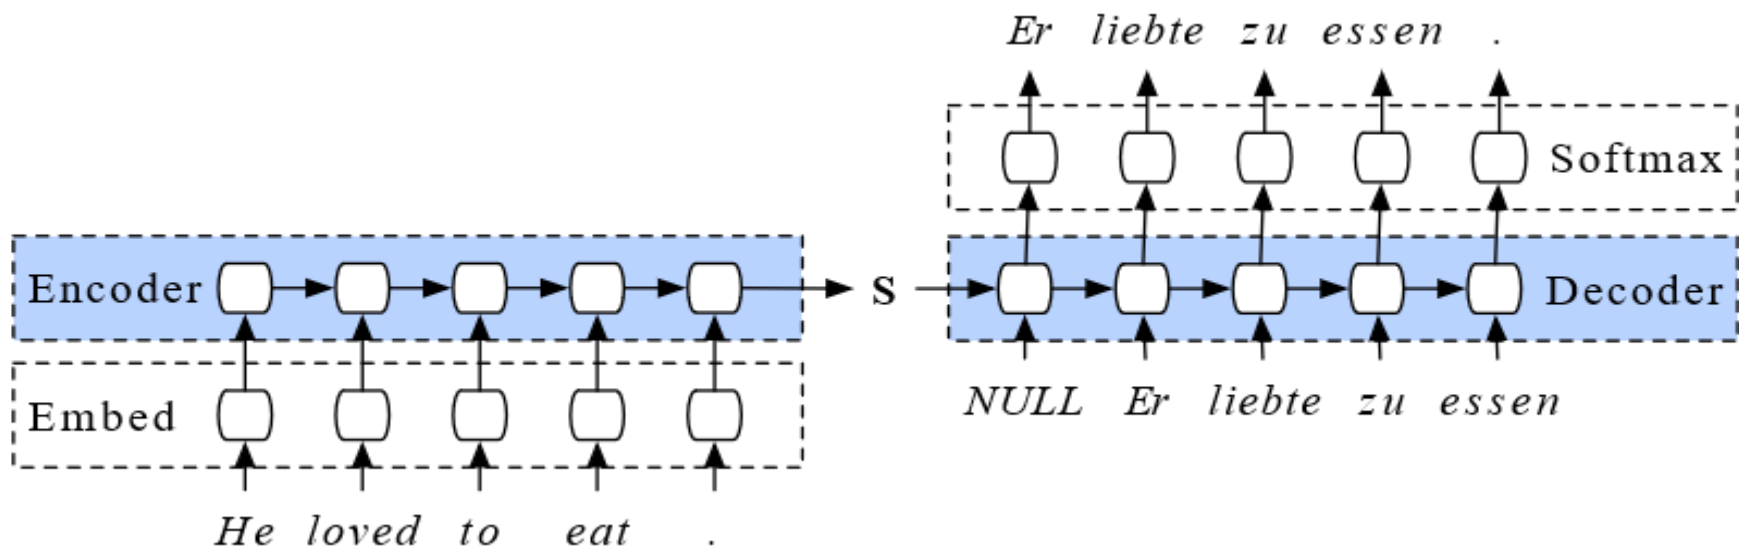
Ulteriori informazioni

Entra in Reverso, è semplice e gratis!

Scopri Ticket Restaurant®, i buoni pasto più spendibili in Italia
Risparmio Costi Aziendali

Neural Machine Translation

<i>Input sentence:</i>	<i>Translation (PBMT):</i>	<i>Translation (GNMT):</i>	<i>Translation (human):</i>
李克強此行將啟動中加總理年度對話機制，與加拿大總理杜魯多舉行兩國總理首次年度對話。	Li Keqiang premier added this line to start the annual dialogue mechanism with the Canadian Prime Minister Trudeau two prime ministers held its first annual session.	Li Keqiang will start the annual dialogue mechanism with Prime Minister Trudeau of Canada and hold the first annual dialogue between the two premiers.	Li Keqiang will initiate the annual dialogue mechanism between premiers of China and Canada during this visit, and hold the first annual dialogue with Premier Trudeau of Canada.



Neural Machine Translation

adottare un vocabolario condiviso è un suggerimento perfetto su come scrivere frasi comprensibili
adopt a shared vocabulary is a perfect suggestion on how to write understandable sentences

un altro suggerimento su come scrivere frasi semplici: evita le negazioni inutili
another suggestion about how to write simple sentences : avoid unnecessary <unk>

quasi 90 persone sono morte per una tempesta tropicale nelle filippine
nearly 90 people died for a tropical storm in the philippines

Figure 3. Some translations from Italian to English by means of the neural model trained by us.

- We have tested the English RCNN model on the same italian SENTIPOLC 2016 test-set translated into English by our neural machine translation model. Results highlight a boost of performance : **78%** of accuracy on the test set versus the **43%** of the Italian trained RCNN model proving our strategy of stacking NMT and RCNN models is successful.

REFERENCES

- Scott, A. J., & Knott, M. (1974).** A cluster analysis method for grouping means in the analysis of variance. *Biometrics*, 507-512.
- Vinyals, O., & Le, Q. (2015).** A neural conversational model. *arXiv preprint arXiv:1506.05869*.
- Bahdanau, D., Cho, K., & Bengio, Y. (2014).** Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*.
- Sutskever I, Hinton GE (2008).** Deep, narrow sigmoid belief networks are universal approximators. *Neural Computation*. 20(11):2629–36. pmid:18533819
- Roux NL, Bengio Y. (2010)** Deep Belief Networks Are Compact Universal Approximators. *Neural Computation*. 22(8):2192–207.
- Amari, S. I., Cichocki, A., & Yang, H. H. (1995, December).** Recurrent neural networks for blind separation of sources. In *Proc. Int. Symp. NOLTA* (pp. 37-42).
- Chung, J., Gulcehre, C., Cho, K., & Bengio, Y. (2014).** Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*.
- Hochreiter S, Schmidhuber J. Long Short-Term Memory (1997).** *Neural Computation*. 1997;9(8):1735–80. pmid:9377276

REFERENCES

Bliemel F. Theil's (1973) Forecast Accuracy Coefficient: A Clarification. *Journal of Marketing Research*. 10(4):444.

Kumar, A., Irsoy, O., Ondruska, P., Iyyer, M., Bradbury, J., Gulrajani, I., ... & Socher, R. (2016, June). Ask me anything: Dynamic memory networks for natural language processing. In *International Conference on Machine Learning* (pp. 1378-1387).

Bow, C., Hughes, B., & Bird, S. (2003, July). Towards a general model of interlinear text. In *Proceedings of EMELD workshop* (pp. 11-13).

Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems* (pp. 3111-3119).

Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. (2014). Generative adversarial nets. In *Advances in neural information processing systems* (pp. 2672-2680).

Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A., & Chen, X. (2016). Improved techniques for training gans. In *Advances in Neural Information Processing Systems* (pp. 2234-2242).

AKNOWLEDGEMENTS



**THANK YOU
FOR YOUR ATTENTION
Francesco Pugliese**